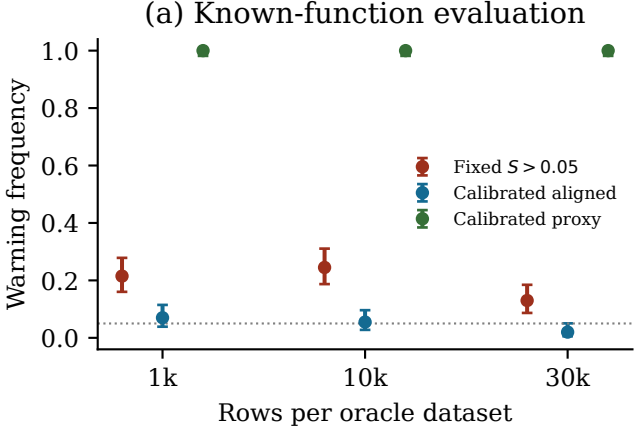
(a) Known-function evaluation
Warming frequency
1.0
0.8
0.6
0.4
0.2
0.0
Fixed $S > 0.05$
Calibrated aligned
Calibrated proxy
1k
10k
30k
Rows per oracle dataset

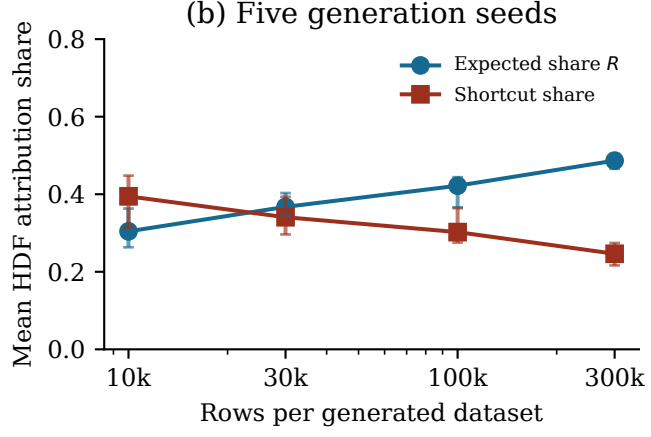
(b) Five generation seeds
Mean HDF attribution share
0.8
0.6
0.4
0.2
0.0
Expected share R
Shortcut share
10k
30k
100k
300k
Rows per generated dataset